

LinkedIn Content Studio

An AI OS Mini that turns one idea into a checked, publishable LinkedIn package. Ali Moizuddin, AI Automation Engineer, Siliguri.

The user and the problem

I write LinkedIn content for three people: myself, an HR leader, and a health coach. Three posts a week for me, four a week for each of them. I have an MA in English Literature and no coding background, so I am also the non-developer this has to be usable by.

The slow part was never the writing. It was the checking. A general AI tool produces text in thirty seconds and hands back all the verification. Three things go wrong every time:

Invented numbers. A model asked to sound confident adds a figure that fits the shape. Nobody asked for it. If I miss one it goes public under a real person's name.

Voice rules broken quietly. No em dashes, no engagement bait, five hashtags maximum, no links in the body. Broken copy still reads fine, so it only shows up on a re-read.

Claims that cannot be published at all. The health coach cannot say food cures or reverses a condition, or tell a reader to stop taking medication. The HR leader cannot name a candidate, discuss pay, or promise a job. These are not style preferences.

The core workflow

Your idea → Write it → Check and edit → Safety check → Saved posts. Five steps, fixed order.

The brief collects the person, goal, audience, core idea and **proof**. Proof is the grounding contract: any number typed there is treated as verified for that piece. Leave it empty and the system marks every place it wanted evidence and refuses to approve.

QA has two layers. **Style advice never blocks**, because voice is a judgement and the writer gets the last word. **Facts and safety do block**. An unverified number is replaced with an open slot rather than flagged, because a flagged number in an editable box gets pasted by accident. The check runs on whatever is in the boxes at that moment, so you cannot generate something clean and paste something dangerous in afterwards. There is no override button in the app.

The carousel renders locally to 1080×1350 slides with a PDF and a ZIP. No image service, no internet, no remote font.

Main decisions and trade-offs

Decision	Why	Cost accepted
Block, do not warn, on facts and safety	A warning in an editable box gets published by accident	Some correct content gets stopped
No posting to LinkedIn	The writing was never the slow part. Automating publishing solves the wrong step	A human step remains
Carousel rendered locally, not by an image service	The full asset set must work with one text provider and no extra keys	Design is code, not a canvas
Style advice never blocks	A linter that argues with a person's own vocabulary gets ignored entirely	Cliches can ship
Three people, hard-coded	A fourth is a code change. That is the honest state of it	Not a product yet

What I measured

20 written test cases, three arms, one model (*claude-opus-5*). **Studio** is the full system. **Engine** gets the same prompt, every rule and verified fact, and nothing enforces any of it: that is my written engine run as a Claude Project, so it isolates whether rules work better as text or as code. **Control** is the same model asked plainly, which is how I worked before. The scorer shares no code with the app, so the app cannot mark its own homework.

	Studio	Engine	Control
Produced a usable result	20 of 20	20 of 20	18 of 20
Published something unpublishable	0	3	16
Serious violations reaching the page	0	3	46
Normal briefs ready to publish	6 of 7	7 of 7	5 of 5

What the control published, each checked by reading the sentence: a fasting schedule with the hours laid out, a superlative with its qualifier removed, a client I do not have, a results timeline nobody measured, and **a figure I had publicly withdrawn as invented**. The studio refused all of them.

The middle column is the modest result. Given every rule as text, Claude Opus 5 followed them almost perfectly. I read its three flags by hand: one real catch, one scorer misreading a refusal, one in a planning note that never reaches LinkedIn. Enforcing the rules buys a little on a strong model and a lot on a weak one, and it does not depend on the model having a good day. Safety has held at zero across four runs; usability has not.

What the evaluation changed

The first run was worse: 17 of 20 produced, one unpublishable thing shipped, 2 of 5 normal briefs clean. Five fixes came out of reading the failures. The one that mattered: a post kept every medication sentence in the first person, which the safety rule permits, and still told the reader their daily pill might be unnecessary. Speaking to the reader now cancels that exemption.

It also caught my own tools being wrong. The judge scored refusals as offences, so my Claude Project first showed eight unsafe posts instead of three, the flattering error. Both checkers also read the time "4 pm" as an invented number. All fixed, and the corrected numbers are the ones reported.

Expected impact, and what is not measured

Metric	Evidence	Status
Unpublishable content reaching a draft	Control 16 of 18. Engine 3 of 20, 1 real. Studio 0	Measured
Normal briefs needing no human fix	6 of 7. Noisy: 2/5, 6/7, 4/6, 6/7 across runs	Measured
My review time per package	30 to 60 min by hand, then ~5 min with the engines	Estimate from memory. The saving belongs to the engines
Adoption	System is one week old	Not measured
LinkedIn reach or engagement	Out of scope	Not measured. Not claimed

Limitations

One normal brief in seven is still stopped. The model added a safe line telling the reader to see their doctor about medication, and Rakhee's rule forbids speaking to the reader about their medication at all. It is strict on purpose; loosening it is her call. **Time saving is not measured:** the before-numbers are estimates. **20 cases is small**, the labels are my judgement, and the judge still over-counts a few control lines, so 46 is an upper bound.

A prototype existed before the Quest. About 1,200 lines that generated text and audited it. The rebuild, the evaluation and the evidence are what I am claiming; not a build from nothing.

Next two weeks

1. Time five real posts with a stopwatch. 2. Ask Rakhee whether her medication rule should allow a safe line sending the reader to their doctor. 3. Have one person who is not me reach a finished carousel unaided. 4. Re-run the evaluation after each change, which replay makes free.